

Joaquin Casares

Staff-Level Platform Engineer – AI Infrastructure, Blockchain, & Developer Experience Leader

Summary: Impact-focused platform architect with 15 years of experience building and rescuing distributed systems. Specializing in AI/ML serving infrastructure (LLMOps), large-scale blockchain networks, and cross-org developer tooling, delivering dramatic improvements in reliability, velocity, and productivity.

- Cut P95 production deployment time by 70% and eliminated a chronic 56% CI/CD failure rate, enabling rapid, reliable releases.
- Automated onboarding across 7 teams, saving 3+ engineer-months per year and slashing new-hire setup from days to minutes.
- Restored stability to an 800+ node Cassandra cluster and operated a 120-node public blockchain network serving 5M+ users.
- Built a staging environment reservation system to eliminate cross-team testing collisions in shared environments.
- Developed diagnostic tooling and knowledge bases still in use after 15 years, providing enduring self-service support resources.

Seeking Staff engineering roles in AI infrastructure, blockchain, or developer experience leadership, with a preferred tech stack including Python/Golang, Kubernetes/Docker, AWS/GCP, Terraform/Helm, and petabyte-scale distributed systems like Cassandra and Ethereum.

Apr 2024 – Present (*1 year*) | **SENIOR SOFTWARE ENGINEER @ WEAVERGRID** Sole engineer focused on **Developer Experience** in the 5-person **Foundations Team** as a resource for **Infrastructure, Kubernetes**, and **CircleCI** support for 30 Engineers.

- Revamp 2+ day onboarding to 40 minutes with setup automation and database access bundles before first 1:1 with manager.
- Eliminate **CircleCI** failure rate of 56%, reduce test runtimes by 73%, and lower P95 deployment duration from 36 to 11 minutes.
- Reduce **Cypress** tests from 8 to 4 minutes and reduce flakes from ~70/week to 2 every 3 weeks via orchestration and tagging.
- Enable **Redis Cluster** Mode and introduce queued deployment strategies to improve critical ingestion and deployment safety.
- Create and automate a cross-service uptime tracking **CircleCI/Cypress** project to alert on reliability gaps across Engineering.
- Independently prioritize critical work for rapid fixes, improving delivery speed and elevating team morale without PM escalation.
- Reduce local **Docker Compose** build times by 70% by optimizing caching and build orchestration to improve developer velocity.

Aug 2023 – Apr 2024 (*8 months*) | **CONSULTANT @ AXONOPS** Produce a performance comparisons and present at the AI/Cassandra Summit, co-lead four training sessions on monitoring **Apache Cassandra**, and publish developer blog posts.

Dec 2021 – Nov 2022 (*1 year*) | **SENIOR SOFTWARE ENGINEER @ AUDIUS** Own CI/CD automation, observability, and release certification for a 119-node network (23 third-party operators) servicing 283M+ requests and 5M+ unique users monthly.

- Automate staging reservations, production deployments, **Ethereum** governance, and **Audius network** releases using **CircleCI**.
- Introduce **Prometheus** metrics and **Grafana** alerts across Slack, Discord, and PagerDuty for **Python** and **TypeScript** codebases.
- Develop first load test against staging using **Node.js** to generate **POA Network**-backed entities proportional to production.

Oct 2018 – Dec 2021 (*3 years*) | **SOFTWARE DEVELOPMENT ENGINEER III @ ADOBE** Decompose the Magento order management monolith into scalable inventory-specific **Go microservices** and support legacy sourcing and config systems.

- Benchmark with **Gatling** scenarios written in **Scala**, collecting results in **Kibana** and **DataDog**, to stress Black Friday readiness.
- Document **GitOps** infrastructure and maintain **Jenkins X**, **Terraform**, **Flux**, **Helm Operator**, **AWS**, **ELK**, and **EKS** environments.
- Define **LDAP**-restricted **Vault** hierarchies and rotate strong secrets consumed by **Kubernetes** via **Sealed** and **Helm Secrets**.
- Migrate from **DataDog** to **Prometheus/Grafana**, **ElasticSearch/Kibana** to **Splunk**, and implement Adobe compliance requests.

Aug 2016 – Oct 2018 (2 years) | **CONSULTANT @ THE LAST PICKLE** Produce runbooks and resolve client issues for **Apache Cassandra** clusters as large as 800+ nodes. TLP was headquartered in New Zealand with 13 geo-distributed employees.

- Automate actionable reports on **Cassandra**, disk, and **Linux** configurations, logs, and **data model** analysis. Still in use in 2024.
- Develop **Grafana** dashboards and alerts while providing critical support for **Cassandra** latency degradation at T-Mobile.
- Create **Docker**-validated runbooks for **Cassandra** maintenance, preventive care, troubleshooting, and **S3/EBS** backup strategies.

Jul 2015 – Aug 2016 (1 year) | **INTERNAL API ENGINEER @ UMBEL** Architect an async **Python** API backed by **Apache Cassandra** and **Dynamo** using **Volatility-Based Decomposition**, **Microservice**, **Event-Driven Design**, and **CQRS** principles.

- Increase **Celery** performance by 3x via asynchronization and **Cassandra** features while reducing **Cassandra** cluster size by 50%.
- Design and implement an **asynchronous data object** for denormalization, client retries, metrics, hydration, and transactions.
- Use **Sentry**, **Jira**, **GitLab**, **Docker Compose**, **Python nose**, and **Slack** for alerting, tracking, development, testing, and discussion.
- Kickstart Umbel Labs innovation through **big data distributed systems**, **blockchain technologies**, and **virtual reality** proposals.

Aug 2014 – Jul 2015 (1 year) | **DEMO ENGINEER @ DATASTAX** Use best practices and common use cases to develop enterprise and community facing demos for **DataStax Enterprise**. Witness 23x growth to 440 geo-distributed employees.

- Develop a full-stack demo, training, and certification platform using **Python**, **Flask**, **Jinja2**, **Bootstrap**, **Docker**, **EC2**, and **Azure**.
- Manage project scope, timeline, and deliverable for demos in cross-team efforts between Marketing, Pre-Sales, and Post-Sales.

Dec 2012 – Jul 2014 (1.7 years) | **SOFTWARE ENGINEER IN TEST @ DATASTAX** Bootstrap **Python**, **Java**, and **C++** client-driver testing harnesses for **Apache Cassandra**. Collaborate with Marketing, Documentation, Training, and Product teams.

- Train 3 engineers and 1 manager to assume ownership of **Java**, **Python**, and **.NET** client-driver testing with cloud-hosted clusters.
- Bootstrap **Python nose** and **Java TestNG** test suites for the **DataStax Drivers for Apache Cassandra** using **Jira** and **Jenkins**.
- Validate, test, and teach Automaton, the internal performance and regression tool run on hardware, **EC2**, and **RightScale**.

Jan 2011 – Jan 2013 (2 years) | **SUPPORT ENGINEER @ DATASTAX** Support **DataStax Enterprise: Apache Hadoop**, **Apache Solr**, and **DataStax Opscenter** on **Apache Cassandra**, in **AWS** and on-prem production environments. Employee #9.

- Primary support contact for 9 months across 16+ products and 200+ organizations; train 5 replacements, including manager.
- Create Knowledge Base that is used daily, as of 2024, and referenced by newly hired and established Support Engineers.
- Contribute to public wikis, design auto-clustering **Amazon Machine Image**, and publish blog posts still driving traffic in 2024.

Aug 2005 – Dec 2010 (5.5 years) | **THE UNIVERSITY OF TEXAS AT AUSTIN**

- Bachelor of Arts **Computer Science**
- Bachelor of Science **Radio-Television-Film**
- Bridging Disciplines Program **Digital Arts & Media**
- New Media Initiative **Advanced Communication Technologies Laboratory**
- Advanced **Digital Media Institute**

joaquin@joaquincares.com | LinkedIn/AngelList/GitHub: joaquincares | joaquincares.com